

Naming the Inexplicable: An Apparatus for Carrying a Yōkai Until It Fades

ANONYMOUS AUTHOR(S)



Fig. 1. **Top:** the five *yōkai* produced by the apparatus whose phenomena have no precedent in the 35,305-entry historical corpus, in order: *Zōshoku-zangyō* (overtime that multiplies), *Kenshū-fujun* (stagnation in a data audit), *Dandenhōkō* (presence at the moment of power loss), *Kōjō-yūei* (faces missing from photographs at twilight), *Tōri-tsuki* (the intrusive image of a colleague). **Bottom:** printed receipts produced by the apparatus during the deployment at YOKAI EXPO 2026 on Shodoshima Island. Each receipt records a single participant's account, the name and narrative the apparatus proposed, and a generated image in folkloric register; the thermal ink fades over months under ambient conditions.

Yōkai is a Japanese category through which strange experiences are placed in a shared cultural vocabulary: a name, a typical place and time, and a short account. Travelers, villagers, ethnographers, woodblock artists, and folklorists have each contributed entries over the centuries. The setting in which an ordinary person contributes a new entry has nearly left contemporary life: the digital archive is consulted for reference rather than used as a substrate for new cultural production. We ask what enters the category when a contemporary individual is given technical means to extend it. We describe an apparatus that performs the *yōkai*-naming act in three to five minutes: a conversational form gathers a personal account; the apparatus retrieves related historical entries; language and image models propose names, narratives, and images in folkloric register, which the participant accepts, modifies, or rejects. The chosen entry is printed on thermal paper that fades over months. Across 79 sessions the apparatus produced 66 entries. We report what was produced and what participants said, and propose that the apparatus is best read not as a machine for adding entries to the folk taxonomy but as a device that compresses a historically communal act into an individual three-minute encounter.

CCS Concepts: • **Applied computing** → **Media arts**; *Digital libraries and archives*; • **Human-centered computing** → Interaction design.

Additional Key Words and Phrases: Computational Folkloristics, Generative AI, Intangible Cultural Heritage, Yōkai, Ephemeral Materialization, Practice-Based Research

1 Introduction

Japanese *yōkai* is a folk category that gives unexplained experience a name and a form. For example, a traveler hears what sounds like someone washing red beans by a stream at night but finds no one; the tradition records this as *Azukiarai*, the bean-washer, attributed regionally to a tanuki (raccoon dog), weasel, or otter [39]. The International

Research Center for Japanese Studies (Nichibunken) has compiled 35,305 such entries in its publicly accessible Yōkai Database [28].

The category has accumulated across periods, from strange phenomena in the Nara-period (710–794) chronicles through Edo-period (1603–1868) woodblock catalogs and Meiji-period (1868–1912) scholarly study to postwar manga and contemporary commercial and online media [10, 13, 16, 23, 27, 31, 40]. Recent scholarship reads *yōkai* as a cultural device that converts inexplicable experience into named cultural objects, indexed to the phenomenon and not to the named entity [24, 39].

The everyday setting in which an ordinary person names a personal unexplained experience has nearly disappeared. Foster identifies that act of naming as central to the taxonomy’s continued growth [11]. Even the wide availability of the Nichibunken corpus has not restored it. The digitized archive is engaged as a reference resource, not as a substrate for participatory production. Some cognate activity continues at the boundary with *kaidan* cultures and regional revival projects, but these are exceptions.

This paper asks what enters the category when a contemporary individual is given technical means to extend it (Fig. 1). We describe an apparatus that performs the *yōkai*-naming act in three to five minutes. A short conversational form gathers the participant’s account. The apparatus retrieves the closest entries from the archive by embedding similarity. A language model proposes names and short narratives in folkloric register; an image model proposes a depiction in one of five inherited pictorial styles. The participant accepts, modifies, or rejects each proposal. The retrieval step anchors each new entry to a specific historical *yōkai* lineage that the participant’s experience extends. The result is printed on thermal paper that fades over months.

2 Related Work

2.1 Yokai studies and naming as productive operation

Three claims from yokai scholarship anchor the apparatus. Komatsu reads *yōkai* as a cultural device that converts inexplicable experience into named cultural objects [24], which supplies the apparatus’s thesis — what is being scaffolded is conversion, not creation. Kagawa attributes the explosive Edo-period expansion of the taxonomy to the productive operation of *naming* itself, traced to a specific cultural condition — the seventeenth-century haikai poet network within which *yōkai* names circulated as *haigon* (haiku words), generating regionally located contributions such as *Ubagabi* [18, 19]. The apparatus targets this same operation, compressing what was historically a distributed network into a single participant encounter. Iikura’s continuing work on contemporary urban legends and *kaidan* culture [21] positions the *yōkai* category as still under active narration in present-day Japan; our finding in Section 4.2 that participants brought workplace, infrastructure, and photographic content to the apparatus is read against this established line.

2.2 Computational engagement with the Nichibunken corpus

The Nichibunken Yōkai Database [28] that the apparatus uses as substrate has been engaged computationally in two prior orientations. Sato and Nakawake apply quantitative methods to the same 35,305-entry corpus to characterize regional distribution and cluster canonical *yōkai* types [20]. Their work establishes that this corpus admits the kind of distributional and categorical comparison we apply in Section 4. Jinnai et al.’s *YokaiEval* uses *yōkai* from Japanese folktales as ground-truth content against which language-model knowledge is evaluated [17]. Both orientations are




Fig. 2. The five stages of the apparatus, from open-ended elicitation through retrieval, naming, visualization, and thermal print. At each stage the AI-generated proposal is visible to the participant, who accepts, modifies, or rejects it. Total processing time is three to five minutes per session.

retrospective — the corpus is the object of analysis. The present apparatus uses the same content prospectively, as substrate against which a participant performs the naming act.

2.3 Constrained generation and ephemeral material form

Two design decisions in the apparatus draw on prior work. The use of generative AI as a constrained creative scaffold, not as an autonomous generator, has been evaluated by Mirowski et al. in language-model-assisted screenwriting with industry professionals [26]. They found the model most useful when it operated as a suggesting collaborator under writer editorial control. We adopt the same configuration in folk-naming, where the apparatus proposes and the participant decides at every stage. The material form of generative output has been studied directly by Ozawa et al., who printed generative outputs through the wet-plate collodion process to ask whether material process can confer on a machine-generated image the standing of a handmade photograph [29]. Thermal printers as a medium for computation that erases itself were preceded in the School for Poetic Computation's installation that lowered thermal output into water as it was produced [33].

What is not yet documented is a system that scaffolds an individual to perform a specific historical folk practice — adding a new entry to a specific historical taxonomy with its naming, narrative, and pictorial conventions — and that materializes the result on a substrate that fades at the timescale of folk transmission. The combination of practice-grounded, individual-scaffolded, taxonomically-constrained, and materially-impermanent is the design space the present work occupies.

3 System Design

Figure 2 summarizes the five computational stages. The stages are grouped below as input, generation, and materialization. At each stage the AI's contribution is a proposal that the participant can accept, modify, or refuse.

The participant describes an experience in free text. The form asks when, where, what was sensed, and what disposition the experience left. These four prompts correspond to features that appear with high frequency in the descriptive text of historical entries, which we coded across a sample of the corpus before designing the form. The participant’s own words are carried through every subsequent stage.

The free-text answers are concatenated into a query, embedded via the Gemini Embedding API (768d), and matched by cosine similarity against pre-computed embeddings of the 35,305 corpus entries [28]. The five closest entries are supplied as context to a language model (Gemini 2.0 Flash), which generates three candidate names and short narratives structured along three patterns Komatsu uses to describe the corpus [24] — phenomenon-descriptive (*Azukiarai*, named for the sound), place-conditional (*Tenjōname*, named for the surface), and sensory-onomatopoeic (*Betobetosan*, named for footsteps). The model is instructed to reference the retrieved entries without reusing existing names. The participant selects, edits, or supplies a different name.

An image generation model (Google’s Nano Banana) then produces a depiction in one of five pictorial registers under which *yōkai* have historically been pictorialized — ink wash (*suiboku-ga*), illustrated scroll (*emaki*), woodblock print (*ukiyo-e*), narrative ink-line drawing (*manga* and *gekiga*), and contemporary digital. The first four are inherited registers; the fifth is included as the visual idiom of the present and is selected when participant accounts describe contemporary settings. Each prompt names conventions internal to the chosen register — Sesshū’s *hatsuboku* line and the eccentric lineage that absorbs *yōkai* imagery [36], the procession structure of the *Hyakki Yagyō Emaki* [35], the catalog grammar of Sekien [31], the postwar *gekiga* tradition — so the model operates within a named tradition. Fixing the categorical frame against drift in model defaults follows Walton’s argument that an image’s representational content depends on the category under which it is perceived [37]. Negative prompts suppress photorealistic, anime, and commercial-mascot aesthetics. Section 4 reports the rate at which the constraint held.

3.1 Materialization

The output is printed on 80mm thermal receipt paper. The fading substrate is the design line along which the apparatus is built. The receipt is small enough to be carried in a wallet, so circulation does not depend on the institutional archive. Thermal ink degrades under ambient light and temperature; the manufacturer specification indicates legibility on the order of months under typical indoor conditions, corresponding approximately to the timescale at which a folk telling might be retold a few times before being forgotten. The fade is structural, not a defect — the entry is allowed to disappear as historical entries did before being recorded, and the apparatus does not commit to remembering on the participant’s behalf. Each session leaves a server-side log of retrieved entries, candidate names, image variants, and the participant’s revisions; the persistent record is the trace of decisions, not the printed artifact.

4 Deployment and Records

4.1 Context and Population

The apparatus was first deployed at YOKAI EXPO 2026 (Fig. 3), a *yōkai* festival held on February 22, 2026 at Fretopia Hall on Shodoshima Island, and continued at successive venues through May 2026. Of 79 initiated sessions, 66 produced a complete *yōkai* (name, narrative, image) and 47 produced a physical receipt. Thirty-seven participants completed both pre- and post-experience surveys. The denominator below is 66 unless otherwise stated. The pre/post category analysis uses the 36 paired surveys with complete responses. Pre-survey means were 3.26 of 5.0 on folklore familiarity (N=72) and 3.51 of 5.0 on prior generative-AI experience (N=69).



Fig. 3. The apparatus in operation at YOKAI EXPO 2026, Fretopia Hall, Tonosho, Shodoshima Island.

4.2 What the Apparatus Produced

Figure 4 shows all 66 outputs at a glance. The 66 generated *yōkai* ranged across the structural types that organize the historical corpus. Phenomenon-descriptive names accounted for the largest share, including entries such as *Tojikoshikage* (a presence seen through closed sliding doors), *Kaidan-nozoki* (a gaze that follows on the stairs), *Shisen-hōhō* (a wandering line of sight), *Yamiyo-no-kodō* (a heartbeat heard in the dark of one's own home), and *Shitsuseirei* (a presence permanently resident in a room). Place-conditional names appeared in entries such as *Mizube-no-me* (an eye at the water's edge) and *Yomichi-mi* (something that looks at the night-road traveler). Sensory and onomatopoeic names appeared in *Chirori*, *Toto-totsu*, *Hiyahiya*, *Don-don-don*, and others. Read together, the corpus of 66 reads as folkloric in register: each entry pairs a place and time with a phenomenon and a named entity, in the form the historical corpus has used.

The distribution of settings, however, differed sharply from the historical corpus (Table 1). Settings that have been productive of *yōkai* across the historical period (home, road, mountain) remained productive in the apparatus output. Settings characteristic of contemporary work and infrastructure (workplace, factory, data audit, electrical infrastructure, event venue, photography) appeared in the apparatus output at rates one to two orders of magnitude above their corpus baseline.

The absence of these contemporary settings from the corpus is structural. A token-level search of all 35,305 corpus entries returned zero occurrences for the Japanese tokens for overtime, audit, power outage, and photographing. The corpus is not blank on every modern setting: school ghost-tales appear (the school category in Table 1 accounts for 0.65%




Fig. 4. Atlas of all 66 *yōkai* produced by the apparatus, in chronological order of generation (left-to-right, top-to-bottom). Three trends are visible across the grid: the intended pictorial registers (ink wash, woodblock, illustrated scroll) dominate the majority of cells; a sub-cluster of contemporary horror-manga and digital-render outputs sits to the right (the partial style permeability discussed below); and the woodblock register is consistently the most legible for entries the apparatus treats as canonical *yōkai*-types. One cell in the bottom-left quadrant is blurred: its output incidentally resembled a well-known copyrighted cartoon character and has been suppressed in the figure while retained in the underlying record.

Table 1. Setting vocabulary distribution. Generated *yōkai* (N=66) versus Nichibunken corpus (N=35,305). Categories defined by token lists; one *yōkai* or entry may match more than one category. Ratio = apparatus rate / corpus rate.

Setting	Gen.%	Corpus %	Ratio
home	43.9	18.13	2.4x
mountain/wilds	9.1	22.32	0.4x
road/path	22.7	8.89	2.6x
water	1.5	17.35	0.1x
school	16.7	0.65	25.7x
workplace/office	7.6	0.03	267.5x
factory	1.5	0.05	29.7x
photography	3.0	0.12	26.1x
event venue	1.5	0.02	66.9x
infrastructure	1.5	0.003	534.9x
data/audit	1.5	0.00	—

of corpus entries), as do some entries on electric lighting and railways. What is absent is the white-collar workplace, the data audit, the digital-imaging device, and the event-venue blackout.

Eight of the 66 generated *yōkai* contained setting vocabulary specific to contemporary life (Fig. 6). The eight are the distinct entries that account for the cross-category counts in Table 1 (one entry may match more than one category, so the count of distinct entries is lower than the sum of per-category counts). Three of the eight combined contemporary

setting with phenomena that have direct historical analogs: *Yūbie* attached the historical phenomenon of missing objects to the office setting; *Chirin* attached visual presence at the periphery of attention to photographs; *Wasureokigaeshi* attached the phenomenon of objects returning to their original place to the workplace. These three extend an existing repertoire into contemporary settings.

The remaining five contain phenomena without precedent in the corpus. *Zōshoku-zangyō* (Multiplying Overtime) is encountered at night in the workplace; work does not diminish but multiplies, as if added to by an unseen force. *Kenshū-fujun* (Inspection Irregularity) appears in the audit of data and records; stagnation enters the records, and rates of progress shift with the affective state of the auditor. *Dandenhōkō* (Power-Outage Wandering) appears at event venues at the moment of an unexpected power loss; a presence wanders the room, drawing warmth from those present. *Kōjō-yūei* (Factory Dusk-Shadow) appears in photographs taken at abandoned factories at twilight; faces are missing from the developed image, and a presence accompanies the absence. *Tōri-tsuki* (Passing Moon) appears during work hours; the image of a difficult colleague passes through the participant's mind, accompanied by a heaviness that arrests motion.

For each of the five, the central phenomenon does not appear among the 35,305 historical entries. The phenomena are not absent from contemporary experience: each was articulated by a visitor as something they had experienced before the encounter with the apparatus. The participant who generated *Kōjō-yūei* described the experience in the post-survey as one in which “through the power of AI, I was able to give material form to phenomena I had only vaguely felt as frightening.” The participant who generated *Tōri-tsuki* wrote: “I was able to generate a *yōkai* from my own experience, and *yōkai* came to feel close to me.” The participant brought the unsettled phenomenon; the apparatus provided a folkloric naming convention through which it could be entered into the taxonomy. The distinguishing content of these five lies in the phenomenon itself, not in the surface form of the name; a language-model bias of the kind documented below under naming convergence does not produce a phenomenon for which the corpus has no precedent.

To characterize the sensibility this cluster expresses, we coded the descriptive text of each generated *yōkai* against a set of recurring tokens grouped along four dimensions (perception, effect, locus, agency) and compared the cluster against the remaining 58 outputs and against a search of the 35,305-entry corpus (Table 2). The cluster shares with the wider apparatus output a high rate of invisibility (75%) and the absence of a named cause (50%). It departs from the wider output and from the corpus on three distinctively contemporary codes: **drain or erosion** as the effect attributed to the entity (warmth, attention, or energy taken without identifiable cause: 25% of the cluster versus 5% of other outputs and 0.5% of the corpus); **multiplication or stagnation** as effect (work that grows, data that does not move: 25% of the cluster, absent from other outputs); and **workplace, institutional-data, infrastructure, and photographic loci**, which together appear in 7 of 8 cluster entries and in 0 of 58 other outputs.

Read together, these codes describe a sensibility distinct from the encounter-on-the-road that the historical corpus more commonly registers. The contemporary cluster names invisible presence located within institutional or labor space, acting through slow drain or unbounded growth, and mediated where it is perceived at all through documents (the photograph, the audit record, the colleague's intrusive image). The historical corpus's most common loci are home (18%), mountain (22%), road (9%), and water (17%); the contemporary cluster reroutes the same naming form onto a built environment that the historical record does not enter. This sensibility was not proposed by the apparatus; the participants brought it. What the apparatus did was give it a place in the taxonomy that has been collecting inexplicable experience for centuries.

A side-by-side comparison of the recurring lexicon makes the contrast direct (Fig. 5). The historical corpus is dominated by tokens for mountain, deity, house, fox, river, village, road, snake, temple, soul, horse, sea, shape-changing,

Table 2. Sensibility coding of generated *yōkai* descriptions. Percentages are the share of entries in which any token from the code appeared. Corpus column reports the same token search across the 35,305-entry Nichibunken corpus.

Code	Cluster	Other gen.	Corpus
perception/invisibility	75%	59%	0.7%
perception/cold	25%	16%	0.7%
effect/drain or erosion	25%	5%	0.5%
effect/multiplication	25%	0%	0.3%
locus/workplace	62%	0%	0.0%
locus/institutional-data	12%	0%	0.1%
locus/infrastructure	12%	0%	0.0%
locus/photographic	25%	0%	0.1%
agency/no-named-cause	50%	43%	1.5%

Historical Nichibunken corpus

35,305 entries reaching from the medieval period through the twentieth century

LANDSCAPE

mountain river sea forest

BUILT settlement

house village road temple shrine

ANIMAL spirit

fox snake horse dog/cat

DEITY / SPIRIT

deity soul monster strange

TIME / WEATHER

night dusk

Generated contemporary cluster

8 entries produced by the apparatus from contemporary participant accounts

WORKPLACE / SCHOOL

workplace school commute overtime work

DOCUMENT / DATA

data audit document

INFRASTRUCTURE

power event venue

PHOTOGRAPHIC

photograph face

AFFECT / EFFECT

cold multiply stagnate erode drain wander mind's-eye colleague

Shared markers across both: invisible presence-felt shape-unseen

Fig. 5. Recurring lexicon of the historical Nichibunken corpus (left, 35,305 entries) and distinctive lexicon of the contemporary cluster (right, 8 entries). Font size scales with the number of entries in which the token appears. The two columns share the markers of invisibility and presence (*kehai*, “shape unseen”), and diverge on locus and effect: the historical corpus is mountain, river, deity, animal-spirit; the contemporary cluster is workplace, school, photograph, infrastructure, drain.

and the umbrella term *kai* (the strange). The contemporary cluster is dominated by workplace, school, photograph, cold, data, power, event, commute, mind’s-eye (*nōri*), and a vocabulary of effect (multiplies, stagnates, erodes, drains, wanders) attached to documents and institutional spaces. The two columns share only the markers of invisibility and presence-without-form. The apparatus matches participant articulation against the left column and returns names that the participants extend into the vocabulary of the right column.

4.3 What Participants Reported

Of the 37 post-experience surveys, 36 contained complete pre- and post-experience category responses. Twenty-four of these 36 (66.7%) ended the encounter in a category different from where they began (Table 3). The post-experience distribution was dominated by the *psychology* category (18 of 36, 50.0%), in which *yōkai* are understood as names given to inarticulate experience. Together with *culture* (7 of 36), the two interpretive categories accounted for 25 of the 36 post-experience responses.




Fig. 6. The five generated *yōkai* whose phenomena have no precedent in the 35,305-entry Nichibunken Yōkai Database. Three (*Zōshoku-zangyō*, *Dandenhōkō*, *Tōri-tsuki*) sit within the apparatus's intended folkloric pictorial registers (woodblock and ink wash). Two (*Kenshū-fujun*, *Kōjō-yūei*) drifted toward a contemporary digital register, an instance of the partial style permeability documented in the failure modes.

Table 3. Pre-post *yōkai* category shift (N=36). Rows: pre-experience category; columns: post-experience category. Bold diagonal: no change.

Pre	char.	cult.	none	psych.	scary	spir.	Σ
character	0	1	0	3	2	1	7
culture	1	5	2	8	1	1	18
none	0	0	1	0	0	0	1
psychology	0	1	1	6	0	0	8
scary	0	0	0	1	0	1	2
Σ	1	7	4	18	3	3	36

All seven visitors who entered with the *character* (commercial-media) category left it after the encounter. Three moved to *psychology*, two to *scary*, one to *culture*, and one to *spiritual*. The pattern was consistent across the operation period: every visitor who arrived holding the commercial-character framing departed from it.

The forced-choice item, asking which of seven thematic descriptions best matched the experience (Table 4), was answered by all 37 post-experience respondents, with up to two selections each, for 61 total selections.

Table 4. Post-experience forced-choice responses (N=37 respondents, up to 2 selections each, 61 total selections). The bar shows the count visually; option G (Other) received zero selections and is omitted.

Survey item	N
A. Enjoyed viewing <i>yōkai</i> characters	12
B. Making <i>yōkai</i> images with AI	18
C. Regional folklore traditions	6
D. Digital technology for tourism	5
E. Visualizing human fears or anxiety	14
F. Transience of memory or tradition	6

The two most-selected items were B (making *yōkai* images with AI, 18 of 37) and E (visualizing human fears and anxiety, 14 of 37). These two responses are not in competition: they describe respectively the technical apparatus and the cultural function that the apparatus was constructed to support.

Among the 23 substantive free-text responses, the recurring thread was an articulation of what the apparatus did, in vocabulary close to the cognitive-mode description developed in Section 1.

- “A work that concretizes and accepts human perceptions and questions.” (10s, *yōkai* enthusiast)
- “Through the power of AI, I was able to give material form to phenomena I had only vaguely felt as frightening.” (20s, general; generated *Kōjō-yūei*)
- “I was able to generate a *yōkai* from my own experience, and *yōkai* came to feel close to me.” (10s, general; generated *Tōri-tsuki*)
- “A work in which one’s own experience easily becomes a *yōkai*.” (40s, other)
- “Visualization of the subconscious.” (20s, general; generated *Okiwasuregami*)
- “A work that visualizes the anxiety and unknown emotions inside oneself when one hears the word *yōkai*.” (40s, general; generated *Tojikoshikage*)
- “A work about contemporary *yōkai*.” (10s, general; generated *Tōri-tsuki*)
- “The concretization of what people today consider *yōkai*-like.” (30s, researcher; generated *Yūgurashioto*)

These eight responses converge on a single description: an experience the participant already had, but had not articulated, is given material form (a name, a narrative, an image) within a folkloric frame. The participants who generated the five *yōkai* without precedent in the corpus describe exactly this transition; the cluster reported in Section 4.2 is not an artifact of the apparatus producing novel content unprompted, but a record of participants bringing contemporary inarticulate phenomena to the apparatus and the apparatus providing the form.

The 1-to-5 measure of intention to investigate folklore or strange tales after the encounter gave a mean of 3.51 (SD 0.99); 21 of 37 respondents (57%) selected 4 or 5.

One participant articulated a limitation of the apparatus directly in the post-experience response: “the concept of cosmic rays from outer space was rendered as image successfully, but since the *aha* experience requires an *aha* moment, no *aha*-like image was generated.” The apparatus produced concrete imagery from concrete description; for experiential structures whose content is the moment of experience itself, the apparatus produced only a static depiction of a setting.

Several visitors were observed folding the printed *yōkai* and placing it in a wallet or notebook, and one participant described the printed result as “the picture given as a souvenir.” The combination of personal generation, brief exchange at the print station, and a printed object that fades over months produced an artifact that participants carried and kept.

Two modes in which the apparatus failed to hold its own constraints appeared consistently across the operation period. **Naming convergence.** Ten of the 66 generated names (15%) share a structural pattern in which the prefix encodes uncertainty about where something is or is not: *Sokohakani*, *Sokohanai* (twice), *Sokokashira*, *Sokokani*, *Sokokana*, *Sokohaka*, *Sokokashiko* (twice), and *Soregadoko*. The convergence may indicate a generative-model bias toward a high-probability semantic ridge, or a substantive feature of the contemporary participant input (uncertainty about the locus of presence, as distinct from the identity of presence). The current data does not separate these readings. **Image style permeability.** The image-generation constraints described in Section 3 prescribed three traditional pictorial registers and suppressed photorealistic and contemporary commercial styles. An inspection of six images sampled across the operation period found that approximately one-third departed from the prescribed registers toward a register familiar from contemporary horror manga. The departures were retained in the record.

5 Discussion

5.1 What the Contemporary Setting Cluster Suggests

The five *yōkai* reported in Section 4.2 (*Zōshoku-zangyō*, *Kenshū-fujun*, *Dandenhōkō*, *Kōjō-yūei*, *Tōri-tsuki*) do not appear in the 35,305-entry corpus. Participants brought the underlying experiences as descriptions of their own lives; what

the apparatus added was a folkloric naming convention through which these experiences entered the vocabulary the corpus has been collecting for centuries.

The phenomenon-based organization of the corpus, as Komatsu has read it [24], implies that what each period names as inexplicable will reflect that period's conditions. The Edo period codified and pictorialized a body of *yōkai* inherited from earlier tradition, centered on the road, the weather, and the waterway. The apparatus, when re-enabling the practice for contemporary participants, produced *yōkai* of the workplace, the data audit, and the infrastructure outage. The cluster is a small empirical instance of what Komatsu's reading would lead one to expect: the cultural device, when operated under different conditions, takes its content from those conditions.

This is not equivalent to producing *yōkai* that will enter living tradition. Whether *Zōshoku-zangyō* or *Tōri-tsuki* circulate beyond the participant who named them, and whether such circulation could happen in the absence of the communal transmission that earlier entries underwent, are open questions that the present records cannot answer. The records support a narrower claim: that the practice, when made available as an individual action under the apparatus, attaches itself to contemporary content. The apparatus commits to a printed object that fades; it does not commit to archiving. This is consistent with the reading offered here: the cultural value addressed is the continuing operation of the practice, not the persistence of any single entry.

The cluster also invites a reading beyond *yōkai* studies. Read with a sociological eye, the entries participants brought to the apparatus point to where contemporary unease accumulates: not the road or the river, but the office at night where work multiplies, the event venue at the moment of an infrastructure failure, the data audit that does not progress, the colleague's image that intrudes on attention during the workday, the abandoned factory photograph in which a face is missing. The historical corpus does not enter these sites; the form they take here describes not an encounter with a localized agent but a slow drain inside institutional and labor space. Contemporary social theory has named related phenomena in other vocabularies (precarity, ambient surveillance, attention capture). The apparatus operates here as a folk probe: it gives names to such phenomena from within the participant's own articulation, in a register inherited from a pre-industrial folk practice. We do not claim the five entries amount to a sociological diagnosis. We claim that the apparatus registers what the participants bring as inexplicable, and what contemporary participants bring concentrates in a specific institutional region of contemporary life.

Read with an art-historical eye, the work positions generative AI at an angle to several recent practices. Sougwen Chung's *Drawing Operations* treats the model as a co-drawing partner whose responses shape the human gesture in real time [7]; Stephanie Dinkins's *Not the Only One* builds a conversational AI grounded in the linguistic and historical particulars of a specific family, refusing the assumption that the model speaks from no one in particular [8]; Lauren Lee McCarthy's *LAUREN* stages the human in the position of a domestic AI assistant [25]. The present apparatus shares with these works the refusal to position the AI as autonomous author, and adds two specific commitments: the model's context is restricted to a single curated heritage corpus, and the artifact is given to the participant on a medium designed to disappear. These choices place the work close to the cultural-probes lineage of Gaver et al. [12] and to the speculative-design tradition [1, 9], in which designed objects function as instruments for cultural inquiry rather than as products. What is added is that the inquiry follows a specific historical folk practice with its own naming and visual conventions, and that the artifact's lifespan matches the timescale of folk transmission.

5.2 From Communal Time to a Single Encounter

The historical formation of a *yōkai* entry is a social process distributed across time. Someone has the experience. They tell it to someone else; the telling varies as it passes through different speakers. A variant takes hold and is shared

across a region. At some point the entry is pictorialized, often long after the first telling. Eventually a compiler records it in writing, and the record enters a corpus. This is not a sequence that one person can perform; it requires retelling, variation, communal acceptance, and historical sedimentation.

The apparatus presented here folds this distributed process into a single encounter that lasts three to five minutes. The participant articulates an experience; the apparatus retrieves related historical entries; the language and image models propose candidates; the participant accepts one; the name and image are printed on receipt paper. The retelling, the regional sharing, and the historical sedimentation that distinguished the historical form do not happen.

The question to ask of this compression is not whether generative AI has revived *yōkai*-making as it existed before. It has not. *Yōkai*-making as it existed before required a communal time that no apparatus can supply. The question to ask is what cultural status the apparatus's output has when an act that was historically completed across decades or longer is completed in a single encounter between an individual and a machine.

The output is not a transmitted *yōkai*: it carries no history of variation, no regional acceptance, no test of survival across retellings. But it is also not merely a machine-generated image. The participant brought the experience; the retrieval step located precedents that the experience could be discussed in terms of; the models proposed a form drawn from those precedents; the participant decided. The result is something that can be read as a *yōkai* for the duration of the encounter — an intersection, temporarily established, of personal experience, historical record-form, and generative proposal.

Read in this way, the apparatus is not a machine for registering new entries in a folk taxonomy. It is a device that lets an individual briefly experience the form of an act that historically belonged to communal time. The articulation, the precedent-matching, the pictorialization, and the inscription that the historical practice distributed across generations are arranged so that a single participant can pass through them in a few minutes. What the participant takes away is not a contribution to the folk tradition; it is the felt structure of the contribution that the tradition once consisted of.

This reading also clarifies why the apparatus combines a historical archive with generative AI. The two are not separable choices that happen to coincide in the design. The archive supplies the historical material that gives the encounter its specificity, so that what the participant performs is read against a particular body of recorded entries rather than against a generic notion of the supernatural. The AI supplies the speed that allows the multi-step process to be traversed within the timeframe of a single visit, so that a participant who arrives without prior preparation can still complete one round of the practice. Either alone would not produce the encounter the apparatus is designed to stage. The archive without the AI would remain a reference resource consulted by experts. The AI without the archive would generate *yōkai*-looking images from its general training without placing the participant in lineage with anything in particular.

The naming convergence reported in Section 4.2 (15% of names sharing a single structural pattern) is one signature of this compression. In the historical form, multiple narrators across regions and generations would have introduced variation that no single language model is positioned to introduce. A version of the apparatus in which entries accumulate across deployments and feed back into later retrieval would partly restore the iterative dimension, making one participant's contribution available as material for another's. The present design does not include this loop, and the present records reflect that limitation.

The participant responses reported in Section 4.3 indicate that at least some participants read the encounter in the terms above. "A work that concretizes and accepts human perceptions and questions"; "a work that visualizes the anxiety and unknown emotions inside oneself when one hears the word *yōkai*"; "the concretization of what people

today consider *yōkai*-like”: these responses describe the encounter as something other than the receipt of a finished cultural object. They describe the participant’s own brief passage through the form of an older practice.

5.3 Limitations

The participant population was self-selected for interest in *yōkai*. The post-experience survey was completed by 37 of the 69 participants who began the pre-survey (54%); the mean folklore familiarity was 3.26 of 5.0. Both figures suggest that the population is not representative of general audiences, although they are less extreme than the festival-only subset would have produced.

The claim that the five *yōkai* reported in Section 4.2 contain phenomena without precedent depends on the corpus and on the token-level search procedure described there. Different searches, or different corpora, could yield different precedents. The claim is one of absence under this corpus and this procedure, not absence in folk tradition.

The image-style constraints are partially permeable, as reported in Section 4.2. The thermal-print fade timing is from manufacturer specification, not from longitudinal observation.

The apparatus depends on commercial language and image models (Gemini 2.0 Flash, Gemini 2.5 Flash Image) whose behavior is subject to revision and whose availability is not under the authors’ control. The records reported here are tied to model behavior at the time of operation. A version of the apparatus that runs on local models would change what kind of artwork the apparatus is and is not yet implemented.

The authors of this paper are not bearers of regional folk tradition. The records and interpretations are open to revision by those who are.

The deployment was conducted in physical exhibition formats; a web deployment would reach beyond exhibition visitors but at the cost of removing the materialization that we treat as central to the encounter, and we are considering this trade-off in subsequent work.

The use of generative AI raises issues of authority and rumor. An apparatus that admits new named entries to a historical taxonomy risks producing material that, taken out of context, could be misread as traditional folk records. We treat this risk seriously: AI generation in this work is a means, not an end. The retrieval step is intended in part as mitigation, by making explicit which historical lineage each new entry extends.

A *yōkai* formed through historical practice is the product of long-running communal accumulation: an experience is told, retold, named, and visually established over decades to centuries. The apparatus borrows the speed of AI to compress the form of this process into a few minutes. It is not intended to substitute apparatus-produced entries for entries refined over communal time; the intent is to allow a visitor to experience the structure of the process — articulating, naming, materializing — with their own input. We separate historiographical fidelity (which historical lineage an entry extends) from artifact content (which is machine-generated), and communicate that separation to participants.

6 Conclusion

Across 79 sessions between February and May 2026, the apparatus produced 66 *yōkai*. Five named experiences without precedent in the 35,305-entry corpus: overtime that multiplies through the night, faces missing from photographs of abandoned factories, a presence at the moment of a power outage, a stagnation in the audit of records, the intrusive image of a difficult colleague during work.

Several participants folded the printed receipt and placed it in a wallet or notebook. The receipt is thermal paper; it will fade. The artwork, on the reading offered here, is not the system that produced the receipt, nor the corpus of

66 entries it produced. It is the act of treating a contemporary inarticulate experience the way Japanese folk tradition treated the bean-washer at the stream: as something to be given a name and a brief written notice and then carried in a pocket until it fades. The five entries without precedent suggest that this act, when performed under contemporary conditions, does not return the practice to its historical content. It allows the practice to do what it has always done, on whatever the present makes available as inexplicable.

The records this paper analyzes are also the first installment of a longer trace. Each session leaves behind the participant's articulation, the retrieved precedents, the candidate names and images, the chosen and the refused, and the final printed entry. Accumulated across deployments, populations, and conditions, such traces would make visible how contemporary inarticulate experience attaches itself to folkloric form across age, region, and occasion. The apparatus has been built so that this future analysis is possible without changing its single-session character: the artifact still leaves with the participant and still fades, but the conditions of its making remain readable. Reading folk practice in this way is the broader research orientation the present paper opens: a folk corpus is no longer only an archive of recorded entries, but also a substrate against which the contemporary act of naming the inexplicable can be observed.

References

- [1] Jeffrey Bardzell and Shaowen Bardzell. 2013. What is "Critical" about Critical Design? In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '13)*. ACM, New York, 3297–3306.
- [2] Genevieve Bell, Mark Blythe, and Phoebe Sengers. 2005. Making by Making Strange: Defamiliarization and the Design of Domestic Technologies. *ACM Transactions on Computer-Human Interaction* 12, 2, 149–173.
- [3] Steve Benford and Gabriella Giannachi. 2012. *Performing Mixed Reality*. MIT Press, Cambridge, MA.
- [4] danah boyd and Kate Crawford. 2012. Critical Questions for Big Data: Provocations for a Cultural, Technological, and Scholarly Phenomenon. *Information, Communication & Society* 15, 5, 662–679.
- [5] Times Higher Education. 2024. How generative AI is helping to educate about Cantonese porcelain. <https://www.timeshighereducation.com>. Accessed 2026.
- [6] Erik Champion. 2015. *Critical Gaming: Interactive History and Virtual Heritage*. Routledge, London.
- [7] Sougwen Chung. 2015–present. *Drawing Operations Unit: Generation 1–3*. Performance and installation series, human-robot drawing. <https://sougwen.com/project/drawing-operations>. Accessed 2026.
- [8] Stephanie Dinkins. 2018–present. *Not the Only One*. Installation: voice-interactive AI built on three generations of African-American family oral history. <https://www.stephaniedinkins.com/ntoo.html>. Accessed 2026.
- [9] Anthony Dunne and Fiona Raby. 2013. *Speculative Everything: Design, Fiction, and Social Dreaming*. MIT Press, Cambridge, MA.
- [10] Michael Dylan Foster. 2009. *Pandemonium and Parade: Japanese Monsters and the Culture of Yōkai*. University of California Press, Berkeley.
- [11] Michael Dylan Foster. 2015. *The Book of Yōkai: Mysterious Creatures of Japanese Folklore*. University of California Press, Berkeley.
- [12] William W. Gaver, Anthony Dunne, and Elena Pacenti. 1999. Design: Cultural Probes. *interactions* 6, 1 (January 1999), 21–29.
- [13] Shigeru Mizuki (original manga, 1965). Toei Animation (TV anime, 1968–present). *GeGeGe no Kitarō*.
- [14] Heritage Trees. 2025. In *SIGGRAPH 2025 Art Gallery*. ACM, New York.
- [15] David J. Hufford. 1982. *The Terror That Comes in the Night: An Experience-Centered Study of Supernatural Assault Traditions*. University of Pennsylvania Press, Philadelphia.
- [16] Enryō Inoue. 1893. *Yōkai-gaku Kōgi* [Lectures on Yōkai-gaku]. Tetsugakkan, Tokyo.
- [17] Yuto Jinnai et al. 2025. Do Large Language Models Know Folktales? A Case Study of Yōkai in Japanese Folktales. *arXiv preprint*. <https://github.com/CyberAgentAILab/YokaiEval>.
- [18] Masanobu Kagawa. 2020. Kimi no Na wa: Naming Yōkai in the Early Modern Period. *Nihon Minzokugaku* [Journal of the Folklore Society of Japan] 302, 1–35.
- [19] Masanobu Kagawa. 2022. *Yōkai o Nazukeru* [Naming the Yōkai]. Iwanami Shoten, Tokyo.
- [20] Kōsuke Sato and Yō Nakawake. 2020. Minkan Denshō no Keiryō Bunseki: *Kaii-Yōkai Denshō Dētabēsu* no Fukan-teki Bunseki [Quantitative Analysis of Folk Tradition: A Statistical Overview of the *Kaii-Yōkai* Tradition Database]. In *Jinmonkon 2020 Proceedings* [Symposium on Humanities and Computer, IPSJ], 30–37.
- [21] Yoshiyuki Iikura. 2025. *Gendai o Ikiru Kaii Yōkai* [*Kaii* and *Yōkai* Living in the Present]. In *Kaii Yōkai-gaku Korekushon* [Collection of *Kaii Yōkai* Studies], Vol. 3, Kazuhiko Komatsu (Ed.). Seikyusha, Tokyo.
- [22] Folger Karsdorp, Marie van der Meer, Enrique Manjavacas, and Mike Kestemont. 2020. Classifying Evaluative Sentiment in International Folktale Variation. *Digital Scholarship in the Humanities* 35, 4, 808–826.

- [23] Kazuhiko Komatsu. 2011. Yōkai to wa nani ka [What are Yōkai?]. In *Yōkai-gaku no Kiso Chishiki* [Foundations of Yōkai-gaku], Kazuhiko Komatsu (Ed.). Kadokawa Gakugei Shuppan, Tokyo, 9–31.
- [24] Kazuhiko Komatsu. 2017. *An Introduction to Yōkai Culture: Monsters, Ghosts, and Outsiders in Japanese History*. Japan Publishing Industry Foundation for Culture (JPIC), Tokyo. Trans. Hiroko Yoda and Matt Alt.
- [25] Lauren Lee McCarthy. 2017. *LAUREN*. Performance work: the artist remotely serves as a custom human intelligent home assistant. <https://lauren-mccarthy.com/LAUREN>. Accessed 2026.
- [26] Piotr Mirowski, Kory W. Mathewson, Jaylen Pittman, and Richard Evans. 2023. Co-Writing Screenplays and Theatre Scripts with Language Models: Evaluation by Industry Professionals. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems (CHI '23)*. ACM, New York.
- [27] Shigeru Mizuki. 2004. *Zusetsu Nippon Yōkai Taizen* [Illustrated Complete Collection of Japanese Yōkai]. Kodansha, Tokyo.
- [28] International Research Center for Japanese Studies (Nichibunken). Yōkai Database (publicly accessible). <https://www.nichibun.ac.jp/youkaidb/>. Accessed 2026.
- [29] Chinatsu Ozawa, Tatsuya Minagawa, and Yoichi Ochiai. 2024. Can AI Generated Ambrotype Chain the Aura of Alternative Process? In *SIGGRAPH Asia 2024 Art Papers*. ACM, New York.
- [30] Noriko T. Reider. 2010. *Japanese Demon Lore: Oni from Ancient Times to the Present*. Utah State University Press, Logan.
- [31] Toriyama Sekien. 1776. *Gazu Hyakki Yagyō* [The Illustrated Night Parade of a Hundred Demons]. Woodblock-printed volumes. Edo (Tokyo).
- [32] Phoebe Sengers, Kirsten Boehner, Shay David, and Joseph “Jofish” Kaye. 2005. Reflective Design. In *Proceedings of the 4th Decennial Conference on Critical Computing: Between Sense and Sensibility (CC '05)*. ACM, New York, 49–58.
- [33] The School for Poetic Computation (SFPC). The Emergence and Decay of Computation. Installation documentation. <https://sfpc.study>. Accessed 2026.
- [34] Hito Steyerl. 2009. In Defense of the Poor Image. *e-flux journal* 10. <https://www.e-flux.com/journal/10/61362/in-defense-of-the-poor-image/>.
- [35] Takako Tanaka. 1994. *Hyakki Yagyō no Mieru Toshi* [The City Where Hyakki Yagyō Is Visible]. Shin'yōsha, Tokyo.
- [36] Nobuo Tsuji. 1970. *Kisō no Keifu* [Lineage of Eccentrics]. Bijutsu Shuppansha, Tokyo. Reissued 2004, Chikuma Shobō.
- [37] Kendall L. Walton. 1970. Categories of Art. *The Philosophical Review* 79, 3, 334–367.
- [38] UNESCO. 2003. Convention for the Safeguarding of the Intangible Cultural Heritage. <https://ich.unesco.org/en/convention>. Accessed 2026.
- [39] Kunio Yanagita. 1956. *Yōkai Dangi* [Discussions on Yōkai]. Shueisha, Tokyo.
- [40] Matthew Meyer. n.d. *yokai.com*: An online database of Japanese ghosts and monsters. <https://yokai.com>. Accessed 2026.
- [41] Level-5 (developer). 2013–present. *Yo-kai Watch* [Video game and media franchise].

AI Tools Disclosure

Large Language Models were used to assist in the research planning and linguistic editing of this manuscript. The generative AI components described in the system (retrieval-augmented generation, image synthesis) are part of the artwork itself. All conceptual design, system implementation, exhibition deployment, observational analysis, and interpretive conclusions were conducted by the human authors, who assume full responsibility for the content.